

Joint convexity of the Umegaki quantum relative entropy

A conventional-mathematics proof, in exact correspondence with the Lean formalization `qRelativeEnt_joint_convexity` (PhysLean)

Prepared for review — PhysLean contribution

2026-07-07

We give a self-contained, conventional-mathematics proof that the Umegaki quantum relative entropy is jointly convex. The argument is presented exactly as it is formalized in Lean 4 in the PhysLean library (theorem `qRelativeEnt_joint_convexity`, commit `f901e353`): every step below corresponds to an identifiable step in the formal proof, and §7 gives the full paper-to-Lean dictionary. The proof route is the Frank–Lieb / Leditzky–Rouzé–Datta variational argument for the sandwiched Rényi trace functional, taken to the limit $\alpha \rightarrow 1^+$; it deliberately avoids differentiability at $\alpha = 1$, using instead a one-sided convergent majorant (Lemma E, the one genuinely new ingredient) that requires only continuity plus two elementary scalar bounds. The purpose of this paper is to let a reviewer confirm the *mathematical* correctness of the contribution in ordinary language, independently of reading Lean, while §7 makes the correspondence to the machine-checked proof precise.

This note renders the Lean proof of `qRelativeEnt_joint_convexity` (QuantumInfo/Entropy/DPI.lean, line 1534 as of commit `f901e353`) as a mathematician would write it on paper. The structure of the argument below follows the Lean proof step for step; Lean identifiers are given in `code font` at the point where they are used. The main theorem delegates its constituent steps to a handful of private supporting lemmas (`mix_M_eq_weighted_sum`, `ker_mix_le`, `sandwichedTraceFunctional_mix_le`, `ofReal_prob_mix_toReal`, `sandwichedRelRentropy_tendsto_qRelative` that sit immediately above it in DPI.lean; §7 gives the full correspondence. Nothing is included that is not in the formal proof, and the two pre-existing ingredients it relies on (joint convexity of the trace functional, and continuity in α) are stated as propositions with their own provenance.

1. Setting and conventions

Throughout, d is a finite index set and states live on $\mathbb{C}^d$.

- A **state** (`MState d`, QuantumInfo/States/Mixed/MState.lean:63) is a Hermitian matrix $\rho \geq 0$ with $\text{Tr } \rho = 1$.
- For $A \geq 0$ and $r \in \mathbb{R}$, the power A^r is defined by functional calculus with $0^r := 0$ for $r \neq 0$ (`HermitianMat.rpow`, QuantumInfo/ForMathlib/HermitianMat/Rpow.lean:45); negative

powers are therefore Moore–Penrose-style powers on the support. Likewise $\log A$ is taken on the support.

- We write $\text{supp } \rho \subseteq \text{supp } \sigma$ for the support (= range) inclusion of PSD matrices; in Lean this is the kernel inclusion $\sigma.\text{M.ker} \leq \rho.\text{M.ker}$.
- Entropies take values in $[0, \infty]$ ($\mathbb{R} \geq 0\infty$), with the conventions $a \cdot \infty = \infty$ for $a \neq 0$ and $0 \cdot \infty = 0$.
- p ranges over $\text{Prob} = \{p \in \mathbb{R} \mid 0 \leq p \leq 1\}$ (`QuantumInfo/ClassicalInfo/Prob.lean:36`). The mixture notation $\mathbf{p} [\rho_1 \leftrightarrow \rho_2]$ (`Mixable.mix`, `Prob.lean:267`) denotes the state whose matrix is $p\rho_1 + (1-p)\rho_2$. In the theorem statement, $1 - \mathbf{p}$ on the right-hand side is truncated subtraction in $\mathbb{R} \geq 0\infty$, which equals the real $1 - p$ since $p \leq 1$.
- $\log$ is the natural logarithm.

2. Definitions

Definition 1 (Umegaki relative entropy; `qRelativeEnt`, notation `D`, `Relative.lean:1497`).

$$D(\rho \parallel \sigma) = \begin{cases} \text{Tr}[\rho(\log \rho - \log \sigma)] & \text{if } \text{supp } \rho \subseteq \text{supp } \sigma, \\ +\infty & \text{otherwise.} \end{cases}$$

(In Lean, $\mathbf{D}(\rho \parallel \sigma)$ is *defined* as $\mathbf{D_1}(\rho \parallel \sigma)$, the $\alpha = 1$ member of the next definition; the identification with the trace formula under the support condition is `qRelativeEnt_ker`, `Relative.lean:1782`.)

Definition 2 (sandwiched Rényi relative entropy; `SandwichedRelRentropy`, notation `D_α`, `Relative.lean:1480`). For $\alpha > 0$, set $\gamma := \frac{1-\alpha}{2\alpha}$. If $\text{supp } \rho \not\subseteq \text{supp } \sigma$, then $\tilde{D}_\alpha(\rho \parallel \sigma) := +\infty$. Otherwise

$$\tilde{D}_\alpha(\rho \parallel \sigma) := \begin{cases} \frac{1}{\alpha-1} \log \text{Tr}[(\sigma^\gamma \rho \sigma^\gamma)^\alpha] & \alpha \neq 1, \\ \text{Tr}[\rho(\log \rho - \log \sigma)] & \alpha = 1, \end{cases}$$

a nonnegative quantity (`sandwichedRelRentropy_nonneg`). (For $\alpha \leq 0$ the Lean definition returns the junk value 0; only $\alpha \geq 1$ is used here.)

Definition 3 (sandwiched trace functional; `sandwichedTraceFunctional`, notation `Q_α`, `DPI.lean:66`). For $\alpha \in \mathbb{R}$ and $\gamma = \frac{1-\alpha}{2\alpha}$,

$$\tilde{Q}_\alpha(\rho \parallel \sigma) := \text{Tr}[(\sigma^\gamma \rho \sigma^\gamma)^\alpha] \in \mathbb{R}.$$

For $0 < \alpha \neq 1$ and $\text{supp } \rho \subseteq \text{supp } \sigma$ one has (`sandwichedRelRentropy_eq_log_traceFunctional`, `DPI.lean:78`)

$$\tilde{D}_\alpha(\rho \parallel \sigma) = \left[\frac{1}{\alpha-1} \log \tilde{Q}_\alpha(\rho \parallel \sigma) \right]^+ \quad (\text{as an element of } [0, \infty], \text{ via } \text{ofReal}).$$

3. Theorem

Theorem (`qRelativeEnt_joint_convexity`, `DPI.lean:1534`). For all states $\rho_1, \rho_2, \sigma_1, \sigma_2$ on $\mathbb{C}^d$ and all $p \in [0, 1]$, writing $\bar{\rho} := p\rho_1 + (1-p)\rho_2$ and $\bar{\sigma} := p\sigma_1 + (1-p)\sigma_2$,

$$D(\bar{\rho} \parallel \bar{\sigma}) \leq p D(\rho_1 \parallel \sigma_1) + (1-p) D(\rho_2 \parallel \sigma_2)$$

as an inequality in $[0, \infty]$.

```

theorem qRelativeEnt_joint_convexity :
  ∀ (ρ₁ ρ₂ σ₁ σ₂ : MState d), ∀ (p : Prob),
    D(p [ρ₁ ↔ ρ₂] || p [σ₁ ↔ σ₂]) ≤ p * D(ρ₁ || σ₁) + (1 - p) * D(ρ₂ || σ₂)

```

4. Ingredients

The proof consumes two substantial pre-existing results (A, B), one elementary matrix lemma (C), and two scalar inequalities from Mathlib (D). Only Lemma E and the assembly in §5 are new in this contribution.

Proposition A (joint convexity of $\tilde{Q}_\alpha$, $\alpha > 1$; `sandwichedTraceFunctional_jointly_convex`, `DPI.lean:568`, pre-existing). Let $\alpha > 1$, let ι be a finite index set, let $w_i \geq 0$ with $\sum_i w_i = 1$, and let $(\rho_i), (\sigma_i)$ be states with $\text{supp } \rho_i \subseteq \text{supp } \sigma_i$ for every i . If $\bar{\rho}, \bar{\sigma}$ are states with $\bar{\rho} = \sum_i w_i \rho_i$ and $\bar{\sigma} = \sum_i w_i \sigma_i$ as matrices, then

$$\tilde{Q}_\alpha(\bar{\rho} \parallel \bar{\sigma}) \leq \sum_i w_i \tilde{Q}_\alpha(\rho_i \parallel \sigma_i).$$

Proof route (already formalized; sketch only). For $\alpha > 1$ and $\text{supp } \rho \subseteq \text{supp } \sigma$ one has the variational formula (`traceFunctional_eq_iSup_f_alpha`, `DPI.lean:434`)

$$\tilde{Q}_\alpha(\rho \parallel \sigma) = \sup_{H \geq 0} f_\alpha(H, \rho, \sigma), \quad f_\alpha(H, \rho, \sigma) = \alpha \text{Tr}[\rho H] - (\alpha - 1) \text{Tr}[(\sigma^{-\gamma} H \sigma^{-\gamma})^{\alpha/(\alpha-1)}]$$

(`f_alpha`, `DPI.lean:174`), the supremum being attained at $\hat{H} = \sigma^\gamma (\sigma^\gamma \rho \sigma^\gamma)^{\alpha-1} \sigma^\gamma$. For fixed $H \geq 0$, f_α is linear in ρ and convex in σ : with $s := -\gamma = \frac{\alpha-1}{2\alpha} \in (0, \frac{1}{2})$ and $q := \frac{\alpha}{\alpha-1} > 1$, the map $\sigma \mapsto \text{Tr}[(\sigma^s H \sigma^s)^q]$ is concave on the PSD cone (`HermitianMat.trace_conj_rpow_concave`, `QuantumInfo/ForMathlib/HermitianMat/LiebConcavity.lean:494`, itself derived from the Lieb–Ando trace theorems formalized in `QuantumInfo/ForMathlib/HayataGroup/TraceInequality/LiebAndoTrace.1` and it enters f_α with the negative coefficient $-(\alpha - 1)$. Hence each $f_\alpha(H, \cdot, \cdot)$ is jointly convex (`f_alpha_jointly_convex`, `DPI.lean:494`), and a pointwise supremum of jointly convex functions is jointly convex (`iSup_f_alpha_jointly_convex`, `DPI.lean:527`). This is the Frank–Lieb variational argument in the form used by Leditzky–Rouzé–Datta; see the literature note. $\square$

Proposition B (continuity in α ; `sandwichedRelRentropy_continuousOn`, `Relative.lean:1769`, pre-existing). For fixed states ρ, σ , the map $\alpha \mapsto \tilde{D}_\alpha(\rho \parallel \sigma)$ is continuous on $(0, \infty)$ (with respect to the order topology on $[0, \infty]$). The consequence used below is:

$$\tilde{D}_\alpha(\rho \parallel \sigma) \longrightarrow \tilde{D}_1(\rho \parallel \sigma) = D(\rho \parallel \sigma) \quad (\alpha \rightarrow 1^+).$$

Note that since $\tilde{D}_1 := D$ by definition, the $\alpha \rightarrow 1$ convergence of the sandwiched Rényi entropy to the Umegaki entropy is precisely the content of this continuity theorem; it is a genuine analytic fact proved earlier in the library, not a definitional triviality.

Lemma C (kernels of weighted sums; `HermitianMat.ker_weighted_sum_le`, `DPI.lean:545`). If $w_i \geq 0$, all $\rho_i, \sigma_i \geq 0$, and $\text{supp } \rho_i \subseteq \text{supp } \sigma_i$ for all i , then $\text{supp}(\sum_i w_i \rho_i) \subseteq \text{supp}(\sum_i w_i \sigma_i)$.

Scalar facts D (Mathlib). (i) $\log x \leq x - 1$ for $x > 0$ (`Real.log_le_sub_one_of_pos`). (ii) $|e^x - 1 - x| \leq x^2$ for $|x| \leq 1$ (`Real.abs_exp_sub_one_sub_id_le`); in particular $e^x - 1 \leq x + x^2$ there.

Also used: strict positivity $\tilde{Q}_\alpha(\rho \parallel \sigma) > 0$ under the support condition (`sandwichedTraceFunctional_pos`, `DPI.lean:100`).

5. The key new lemma: a convergent majorant for the difference quotient

Lemma E (`sandwichedTraceFunctional_sub_one_div_eventually_le`, `DPI.lean:1432, new`). Let ρ, σ be states with $\text{supp } \rho \subseteq \text{supp } \sigma$ (so $D(\rho\|\sigma) < \infty$). Then there is a function $u : \mathbb{R} \rightarrow \mathbb{R}$ with

$$u(\alpha) \xrightarrow{\alpha \rightarrow 1^+} D(\rho\|\sigma) \quad \text{and} \quad \frac{\tilde{Q}_\alpha(\rho\|\sigma) - 1}{\alpha - 1} \leq u(\alpha) \quad \text{for all } \alpha > 1 \text{ sufficiently close to } 1.$$

Proof. For $\alpha > 1$ put

$$r(\alpha) := \frac{\log \tilde{Q}_\alpha(\rho\|\sigma)}{\alpha - 1}, \quad x(\alpha) := (\alpha - 1)r(\alpha) = \log \tilde{Q}_\alpha(\rho\|\sigma),$$

and define $u(\alpha) := r(\alpha) + x(\alpha)r(\alpha) = r(\alpha) + (\alpha - 1)r(\alpha)^2$.

1. Under the support condition, $r(\alpha) \geq 0$ and $\tilde{D}_\alpha(\rho\|\sigma) = [r(\alpha)]^+$ (Definition 3 with `sandwichedRelEntropy_nonneg`), and $D(\rho\|\sigma) \neq \infty$ (Lean: `h_ne`, from `qRelativeEnt_ne_top_iff.mpr hker`; `Relative.lean:1789`). By Proposition B, packaged as `sandwichedRelEntropy_tendsto_qRelative` (`DPI.lean:1423`), $r(\alpha) \rightarrow D(\rho\|\sigma)$ (as real numbers) as $\alpha \rightarrow 1^+$ (Lean: `h_r_tendsto`, passing through `ENNReal.tendsto_toReal`).
2. Hence $x(\alpha) = (\alpha - 1)r(\alpha) \rightarrow 0$, so eventually $|x(\alpha)| \leq 1$ (Lean: `h_eps`, `h_small`).
3. $\tilde{Q}_\alpha(\rho\|\sigma) > 0$ (positivity above), so $\tilde{Q}_\alpha(\rho\|\sigma) = e^{x(\alpha)}$. By D(ii), for the eventual α of step 2, $\tilde{Q}_\alpha - 1 \leq x(\alpha) + x(\alpha)^2$; dividing by $\alpha - 1 > 0$,

$$\frac{\tilde{Q}_\alpha(\rho\|\sigma) - 1}{\alpha - 1} \leq r(\alpha) + x(\alpha)r(\alpha) = u(\alpha).$$

4. By steps 1–2, $u(\alpha) \rightarrow D(\rho\|\sigma) + 0 \cdot D(\rho\|\sigma) = D(\rho\|\sigma)$. ■

Heuristic remark (not part of the formal proof). Since $\tilde{Q}_1(\rho\|\sigma) = \text{Tr } \rho = 1$, the quantity $(\tilde{Q}_\alpha - 1)/(\alpha - 1)$ is a difference quotient whose two-sided limit “should” be $\frac{d}{d\alpha} \tilde{Q}_\alpha|_{\alpha=1} = D(\rho\|\sigma)$. The formal proof deliberately avoids establishing differentiability: a one-sided *eventual upper bound by a convergent majorant* suffices for the limit argument, and that follows from continuity of $\tilde{D}_\alpha$ alone, via the two elementary scalar bounds.

6. Proof of the theorem

Fix $\rho_1, \rho_2, \sigma_1, \sigma_2$ and p . The proof splits exactly as the Lean source does.

Step 0 — degenerate weights. If $p = 0$, then $\bar{\rho} = \rho_2$, $\bar{\sigma} = \sigma_2$ (in Lean via `Mixable.mix_zero`, `Prob.lean:279`), and the right-hand side is $0 \cdot D(\rho_1\|\sigma_1) + 1 \cdot D(\rho_2\|\sigma_2) = D(\rho_2\|\sigma_2)$ — note this uses $0 \cdot \infty = 0$, so the case is genuine, not vacuous. Equality holds. If $p = 1$, symmetrically $\bar{\rho} = \rho_1$, $\bar{\sigma} = \sigma_1$ via the companion mixture identity `Mixable.mix_one` (`Prob.lean:284`, a new `@[simp]` lemma added by this contribution), and the right side is $D(\rho_1\|\sigma_1)$. Both degenerate cases now close by `simp` alone, the two `@[simp]` mixture lemmas discharging the mixture identity. Henceforth $0 < p < 1$ (Lean: `hp0'`, `hp1'`).

Step 1 — infinite right-hand side. Suppose $\text{supp } \rho_1 \not\subseteq \text{supp } \sigma_1$ (Lean: `hker_1` fails). Then $D(\rho_1\|\sigma_1) = \infty$ by Definition 1 — in Lean via `qRelativeEnt_eq_top_iff.mpr hker_1` (`Relative.lean:1795`), one of the two new finiteness characterizations — and since $p \neq 0$,

$p \cdot \infty = \infty$, so the right-hand side is ∞ and the inequality is trivial (`le_top`). The case $\text{supp } \rho_2 \not\subseteq \text{supp } \sigma_2$ is identical, using $1 - p \neq 0$. So we may assume both support conditions

$$\text{supp } \rho_i \subseteq \text{supp } \sigma_i \quad (i = 1, 2),$$

whence $D(\rho_1 \| \sigma_1), D(\rho_2 \| \sigma_2) < \infty$; in Lean this finiteness is supplied where needed by the dual lemma `qRelativeEnt_ne_top_iff.mpr hker_i` (`Relative.lean:1789`).

Step 2 — setup for the main case. Write the mixtures as weighted sums over $\iota = \{1, 2\}$ (Lean: `Fin 2`) with weight vector $w = (p, 1 - p)$:

$$\bar{\rho} = \sum_i w_i \rho_i, \quad \bar{\sigma} = \sum_i w_i \sigma_i$$

as matrix identities. In Lean this weighted-sum form is packaged, once and for all, as the lemma `mix_M_eq_weighted_sum` (`DPI.lean:1485`), which unfolds `Mixable.mix` into $\sum_i w_i (\cdot)_i$ with $w = (p, 1 - p)$. By Lemma C, $\text{supp } \bar{\rho} \subseteq \text{supp } \bar{\sigma}$; the weight nonnegativity, the weight sum $\sum_i w_i = 1$, and the per-component support conditions are all discharged inside the dedicated lemma `ker_mix_le` (`DPI.lean:1493`), so the theorem simply sets `hker_mix := ker_mix_le p hker_1 hker_2`. Apply Lemma E to each pair to obtain majorants u_1, u_2 with $u_i(\alpha) \rightarrow D(\rho_i \| \sigma_i)$ as $\alpha \rightarrow 1^+$ and, eventually, $(\tilde{Q}_\alpha(\rho_i \| \sigma_i) - 1)/(\alpha - 1) \leq u_i(\alpha)$.

Step 3 — pointwise bound for $\alpha > 1$ (Lean: `h_ev`). The binary case of Proposition A, specialized to a `Mixable` mixture, is packaged as the lemma `sandwichedTraceFunctional_mix_le` (`DPI.lean:1504`); `h_ev` invokes it as `sandwichedTraceFunctional_mix_le hα p hker_1 hker_2`. For $\alpha > 1$ in the eventual range of both majorants, using $\tilde{Q}_\alpha(\bar{\rho} \| \bar{\sigma}) > 0$ (positivity, with `hker_mix`) and Definition 3:

$$\tilde{D}_\alpha(\bar{\rho} \| \bar{\sigma}) = \left\lceil \frac{\log \tilde{Q}_\alpha(\bar{\rho} \| \bar{\sigma})}{\alpha - 1} \right\rceil^+,$$

and we bound the bracketed real number:

$$\begin{aligned} \frac{\log \tilde{Q}_\alpha(\bar{\rho} \| \bar{\sigma})}{\alpha - 1} &\leq \frac{\tilde{Q}_\alpha(\bar{\rho} \| \bar{\sigma}) - 1}{\alpha - 1} && [\text{D(i): } \log x \leq x - 1] \\ &\leq \frac{p \tilde{Q}_\alpha(\rho_1 \| \sigma_1) + (1-p) \tilde{Q}_\alpha(\rho_2 \| \sigma_2) - 1}{\alpha - 1} && [\text{Prop. A at } \iota = \{1, 2\}: \text{ sandwichedTraceFunctional_mix_le}] \\ &= p \frac{\tilde{Q}_\alpha(\rho_1 \| \sigma_1) - 1}{\alpha - 1} + (1-p) \frac{\tilde{Q}_\alpha(\rho_2 \| \sigma_2) - 1}{\alpha - 1} && [\text{since } p + (1-p) = 1] \\ &\leq p u_1(\alpha) + (1-p) u_2(\alpha) && [\text{Lemma E}]. \end{aligned}$$

Hence, eventually as $\alpha \rightarrow 1^+$,

$$\tilde{D}_\alpha(\bar{\rho} \| \bar{\sigma}) \leq \lceil p u_1(\alpha) + (1-p) u_2(\alpha) \rceil^+.$$

Step 4 — pass to the limit $\alpha \rightarrow 1^+$. By Proposition B (applied to the mixture pair through `sandwichedRelEntropy_tendsto_qRelativeEnt`, `DPI.lean:1423`), the left side tends to $D(\bar{\rho} \| \bar{\sigma})$ (Lean: `h_lhs`). By Lemma E and continuity of $t \mapsto [t]^+$, the right side tends to

$$\lceil p D(\rho_1 \| \sigma_1) + (1-p) D(\rho_2 \| \sigma_2) \rceil^+ = p D(\rho_1 \| \sigma_1) + (1-p) D(\rho_2 \| \sigma_2)$$

(Lean: `h_rhs`; the identification of the $[\cdot]^+$ of the real convex combination with the $[0, \infty]$ -valued convex combination is packaged as the lemma `ofReal_prob_mix_toReal` (`DPI.lean:1519`), which consumes the finiteness from Step 1 supplied by `qRelativeEnt_ne_top_iff`). Since the inequality of Step 3 holds eventually along the (nontrivial) filter of right-neighbourhoods of 1, the inequality passes to the limits (`le_of_tendsto_of_tendsto`):

$$D(\bar{\rho} \| \bar{\sigma}) \leq p D(\rho_1 \| \sigma_1) + (1-p) D(\rho_2 \| \sigma_2). \quad \blacksquare$$

7. Lean–paper dictionary

Paper object / step	Lean identifier	Location (commit f901e353)
Umegaki D	<code>qRelativeEnt / D</code>	<code>Relative.lean:1497</code>
Sandwiched $\tilde{D}_\alpha$	<code>SandwichedRelEntropy / D_α</code>	<code>Relative.lean:1480</code>
Trace functional $\tilde{Q}_\alpha$	<code>sandwichedTraceFunctional / Q_α</code>	<code>DPI.lean:66</code>
$\tilde{D}_\alpha = [\log \tilde{Q}_\alpha / (\alpha - 1)]^+$	<code>sandwichedRelEntropy_eq_logDPofRealFunctional</code>	<code>Relative.lean:568</code>
Proposition A	<code>sandwichedTraceFunctional_joinDPofRealFunctional</code>	<code>Relative.lean:568</code>
— variational formula	<code>traceFunctional_eq_iSup_f_alpha, H_hat</code>	<code>PIA.lean:434, 174, 180</code>
— concavity in σ	<code>HermitianMat.trace_conj_rpow_LiebConcavity</code>	<code>lean:494</code>
Proposition B	<code>sandwichedRelEntropy.continuous</code>	<code>Relative.lean:1769</code>
— packaged as $\alpha \rightarrow 1^+$ convergence	<code>sandwichedRelEntropy_tendsto_qRelativeEnt</code>	<code>DPI.lean:428</code>
Lemma C	<code>HermitianMat.ker_weighted_sum</code>	<code>DPI.lean:545</code>
D(i)	<code>Real.log_le_sub_one_of_pos</code>	<code>Mathlib</code>
D(ii)	<code>Real.abs_exp_sub_one_sub_id</code>	<code>Mathlib</code>
$\tilde{Q}_\alpha > 0$	<code>sandwichedTraceFunctional_pos</code>	<code>DPI.lean:100</code>
Lemma E	<code>sandwichedTraceFunctional_subDPofReal</code>	<code>DPI.lean:1432</code>
$\mathbf{D} \neq \mathbf{T} \iff \text{supp-inclusion}$	<code>qRelativeEnt_ne_top_iff</code>	<code>Relative.lean:1789</code>
(Step 1 finiteness)		
$\mathbf{D} = \mathbf{T} \iff \text{supp-inclusion fails}$	<code>qRelativeEnt_eq_top_iff</code>	<code>Relative.lean:1795</code>
(Step 1 ∞ -branch)		
Mixture as $\sum_i w_i(\cdot)_i$ (Step 2 matrix identity)	<code>mix_M_eq_weighted_sum</code>	<code>DPI.lean:1485</code>
$\text{supp } \tilde{\rho} \subseteq \text{supp } \tilde{\sigma}$ (Step 2, <code>hker_mix</code>)	<code>ker_mix_le</code>	<code>DPI.lean:1493</code>
Binary $\tilde{Q}_\alpha$ convexity (Step 3) ofReal of the convex combination (Step 4 identity)	<code>sandwichedTraceFunctional_mixDPofReal_prob_mix_toReal</code>	<code>DPI.lean:1504</code> <code>DPI.lean:1519</code>
$p = 0$ / $p = 1$ degenerate mixtures (Step 0)	<code>Mixable.mix_zero / Mixable.mix_one</code>	<code>Prob.lean:279, 284</code>
Steps 0–4	<code>qRelativeEnt_joint_convexity</code>	<code>DPI.lean:1534</code>
Step 3	<code>h_ev</code> (calling <code>sandwichedTraceFunctional_mixDPofReal</code>)	<code>DPI.lean:1534</code>
Step 4	<code>h_lhs, h_rhs, le_of_tendsto_of_tendsto</code>	<code>DPI.lean:1534</code>

8. Remarks for the reviewer

- **What is new.** Lemma E and the theorem itself, together with the five private supporting lemmas that carry the individual proof steps (`sandwichedRelEntropy_tendsto_qRelativeEnt`, `mix_M_eq_weighted_sum`, `ker_mix_le`, `sandwichedTraceFunctional_mix_le`, `ofReal_prob_mix_toReal`), account for the +191 lines added to `DPI.lean` (§5–§6). In `Relative.lean` the previously

`@[sorryful]` statement was removed and replaced by the two finiteness characterizations `qRelativeEnt_ne_top_iff / qRelativeEnt_eq_top_iff` (`Relative.lean:1789, 1795; net +11 / -15`), and `Prob.lean` gains one `@[simp]` lemma, `Mixable.mix_one` (`Prob.lean:284, +5`). Propositions A and B and all deeper inputs (Lieb–Ando trace theorems) were already in the library. The theorem was placed in `DPI.lean` rather than `Relative.lean` because the $\tilde{Q}_\alpha$ machinery lives there and importing it from `Relative.lean` would create a cycle.

- **No differentiability is needed.** The classical way to see $\tilde{D}_\alpha \rightarrow D$ and to transfer convexity is to differentiate $\alpha \mapsto \tilde{Q}_\alpha$ at $\alpha = 1$; the formal proof replaces this by the one-sided majorant of Lemma E, which needs only continuity (Proposition B) plus the scalar bounds D . This is the main proof-engineering choice.
- **The statement is the binary form.** Finite mixtures of n states follow by induction in the usual way, but only the n -ary form of Proposition A (arbitrary finite `Fintype ι`) is formalized; the **D**-level statement is for two pairs, matching the pre-existing `sorryful` statement it closes.
- **Axioms.** The proof is `sorry`-free; `#print axioms qRelativeEnt_joint_convexity` reports only `propext`, `Classical.choice`, `Quot.sound` (recorded in the commit message of `f901e353`).